

OpenMP — $N = 1024^3$ — speedup over each machine's own single-GPU run. Higher is better;
 1.0 = the extra GPUs bought nothing. [1-GPU wall time: H100 34.7 s · Max-1550 24.6 s · MI250X 73 s]

