

OpenMP — $N = 1152^3$ — speedup over each machine's own single-GPU run. Higher is better;
 1.0 = the extra GPUs bought nothing. [1-GPU wall time: H100 49.6 s · PVC 1550 35.8 s · MI250X 104 s]

