

OpenMP — $N = 1280^3$ — speedup over each machine's own single-GPU run. Higher is better;
 1.0 = the extra GPUs bought nothing. [1-GPU wall time: H100 65.7 s · Max-1550 48.9 s · MI250X 142 s]

